Introduction to Monte Carlo in Finance

Giovanni Della Lunga

WORKSHOP IN QUANTITATIVE FINANCE

Bologna - May 18-19, 2017

Outline

What is Monte Carlo?

- From a quite general point of view (not a really precise one, actually) with the term Monte Carlo usually one means a numerical technique which makes use of random numbers for solving a problem.

- For the moment we assume that you can understand, at least intuitively, what a random number is.

- Later we will return to the definition of a random number, and, as we shall see, this will lead to absolutely not trivial issues.

- Let's start immediately with some practical examples (we'll try to give a more formal definition later).

What is Monte Carlo?

- Let's consider two problems apparently very different in nature;

- The first one is of probabilistic nature: the assessment of the premium for an European Option on a stock that does not pay dividends;

- The second one is an issue of purely deterministic nature: the determination of the area enclosed by a plane figure, such as a circle.

What is Monte Carlo?

- Let's start with the first problem. The pricing of an option is usually dealt with in the context of so-called risk-neutral valuation.

- Indicating with $f[S(T)]$ where S is the value of the underlying asset, the value of the option at maturity T, the value today, $f[S(t)]$, is given by

$$f(S(t)) = \mathbb{E}^{\mathbb{Q}}\left[P(t, T)f[S(T)]\right]$$

- $\mathbb{E}^{\mathbb{Q}}$ being the risk-neutral expectation value and $P(t, T)$ the discount function between t and T.

- Let's assume, for simplicity, to know with certainty the value of the discount function so the problem can be put in the form

$$f(S(t)) = e^{-r(T-t)}\mathbb{E}^{\mathbb{Q}}\left[f[S(T)]\right]$$

What is Monte Carlo?

- The formulation of the problem makes clear its inherently probabilistic nature.

- The application of the Monte Carlo method in this case is reduced essentially to the generation of a sufficiently high number of estimates of $f[S(T)]$ from which to extract the average value.

⚠ To this end it is necessary first to introduce a hypothesis on how the underlying stock price evolves over time;

What is Monte Carlo?

- Let's suppose for example that the asset price follows a geometric Brownian motion, according to this hypothesis the rate of change of the price in a range of infinitesimal time is described by

$$dS_t = rS_t dt + S_t \sigma dZ_t$$

  where r is the risk free rate, σ is the volatility of S returns and dw is a brownian motion;

- A discrete version, which can easily be simulated is given by the difference equation

$$S_t = S_{t-\Delta t} exp \left[\left(r - \frac{1}{2}\sigma^2 \right) \Delta t + \sigma \sqrt{\Delta t} z_t \right]$$

  for times $t \in (\Delta t, 2\Delta t, \ldots, T)$ and z_t being standard normally distributed random numbers;

What is Monte Carlo?

- Once we have the simulated value of the underlying at time T, we are able to derive the value of the option at the same date;
- Assuming for example that the option is a CALL we simply write

$$f[S(T)] = \max(S(T) - K, 0)$$

  where K is the strike price.

- By repeating the above procedure a very large number of times we are able to obtain a distribution of values for $f[S(T)]$ from which it is possible to extract the expectation value ...

What is Monte Carlo?

Interlude – Let's start coding...

Out toolbox: Jupyter, Python, R

- Python, R

- The Python world developed the IPython notebook system.

- Notebooks allow you to write text, but you insert code blocks as "cells" into the notebook.

- A notebook is interactive, so you can execute the code in the cell directly!

- Recently the Notebook idea took a much enhanced vision and scope, to explicitly allow languages other than Python to run inside the cells.

- Thus the Jupyter Notebook was born, a project initially aimed at Julia, Python and R (Ju-Pyt-e-R). But in reality many other languages are supported in Jupyter.
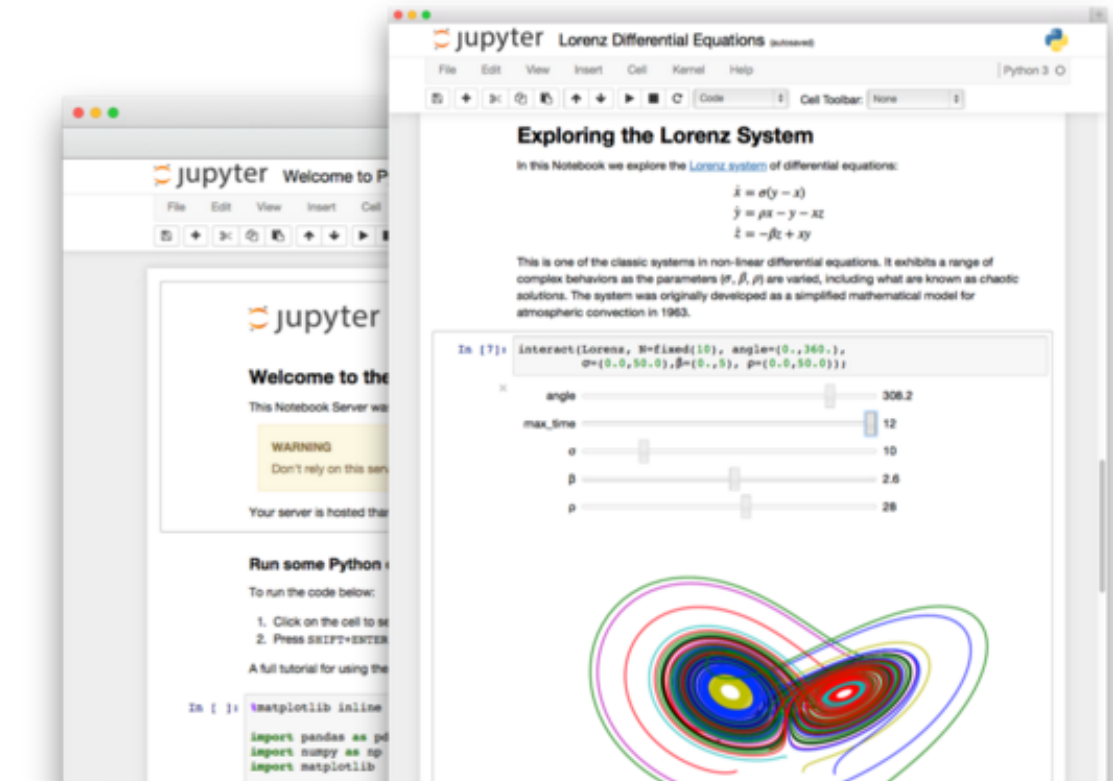
Out toolbox: Jupyter, Python, R

Jupyter Notebook

The Jupyter Notebook is a web application that allows you to create and share documents that contain live code, equations, visualizations and explanatory text. Uses include: data cleaning and transformation, numerical simulation, statistical modeling, machine learning and much more.

Out toolbox: Jupyter, Python, R

Out toolbox: Jupyter, Python, R

- **The IRKernel**

- To enable support of a new language means that somebody has to write a "kernel".

- The kernel for R is called IRKernel (available at github).

- **How do you use Jupyter?**

- Once Jupyter is up and running, you interact with it on a web page.

Out toolbox: Jupyter, Python, R

- **Benefits of using Jupyter**

- Jupyter was designed to enable sharing of notebooks with other people. The idea is that you can write some code, mix some text with the code, and publish this as a notebook. In the notebook they can see the code as well as the actual results of running the code.

- This is a nice way of sharing little experimental snippets, but also to publish more detailed reports with explanations and full code sets. Of course, a variety of web services allows you to post just code snippets (e.g. gist). What makes Jupyter different is that the service will actually render the code output.

- One interesting benefit of using Jupyter is that Github magically renders notebooks. See for example, the github Notebook gallery.

Notebook

- **GitHub :**   polyhedron-gdl;
- **Notebooks :**   mcs_1;

What is Monte Carlo?

Regardless the many different definitions, Montecarlo methods share a common procedural pattern;

1. Define a domain of possible inputs;
2. Generate inputs randomly from a probability distribution over the domain;
3. Perform a deterministic computation on the inputs;
4. Aggregate the results;

This is particularly evident in the next example...

What is Monte Carlo?

$n = 3000 \, (\pi \approx 3.16667)$

- Let's now consider the second problem;

- Take a circle inscribed in a unit square. Given that the circle and the square have a ratio of areas that is $\pi/4$, the value of π can be approximated using a Monte Carlo method:

What is Monte Carlo?

$n = 12000\ (\pi \approx 3.14633)$

1. Draw a square on the ground, then inscribe a circle within it.

2. Uniformly scatter some objects of uniform size (grains of rice or sand) over the square.

3. Count the number of objects inside the circle and the total number of objects.

4. The ratio of the two counts is an estimate of the ratio of the two areas, which is $\pi/4$. Multiply the result by 4 to estimate π.

What is Monte Carlo?

$n = 21000 \ (\pi \approx 3.14286)$

- In this procedure the domain of inputs is the square that circumscribes our circle.

- We generate random inputs by scattering grains over the square then perform a computation on each input (test whether it falls within the circle).

- Finally, we aggregate the results to obtain our final result, the approximation of π.

What is Monte Carlo?

$n = 21000 \ (\pi \approx 3.14286)$

- There are two important points to consider here...

- First, if the grains are not uniformly distributed, then our approximation will be a poor one;

- Second, there should be a large number of inputs. The approximation is generally very poor if only a few grains are randomly dropped into the whole square. On average the approximation improves (slowly!) as more grains are dropped.

Notebook

- **GitHub :** polyhedron-gdl;
- **Notebooks :**   mcs_2;

Monte Carlo is Integration!

- There is a formal connection between the use of the Monte Carlo method and the concept of integration of a function.
- First of all we observe how the problems discussed in the previous paragraph can be attributed both to the calculation of integrals.
- The case related to the area of the circle is evident
- The price of an option as we have seen is nothing more than the discounted value of the expectation value of the price at maturity, the underlying risk factor (the stock price) is distributed according to a log-normal distribution, therefore, we have (for the CALL case):

$$C(t,S) = e^{-r(T-t)} \int_{0}^{+\infty} max[S(T) - K]\, \phi[S(T)]\, dS(T)$$

Monte Carlo is Integration!

- More in general we can state that each extraction of a sample of random numbers can be used as an estimator of an integral.
- As an example consider the case relating to the integration of a function of a real variable;
- by a suitable change of variable, we can always bring us back to the simplest case in which the integration interval is between 0 and 1:

$$I = \int_0^1 f(x)dx$$

Monte Carlo is Integration!

- The key point of our argument is to recognize that the expression written above is also the expectation value of the function f at values of a random variable uniformly distributed in the range $[0, 1]$.

- It becomes possible to estimate the value of our integral using an arithmetic mean of n values of $f(U_i)$ where each U_i is a sample from a uniform distribution in $[0, 1]$.

- In other words we can say that the quantity

$$\tilde{I}_n = \frac{1}{n} \sum_{i=1}^{n} f(U_i)$$

is an unbiased estimator of I.

Monte Carlo is Integration!

- The variance of the estimator is

$$var\left(\tilde{I}_n\right) = \frac{var(f(U_i)}{n}$$

- the mean square error of the estimator, which can be interpreted as the mean square error of the Monte Carlo simulation, decreases with increasing n.

- This result is completely independent of the dimensionality of the problem.

- It's this last characteristic that makes attractive the Monte Carlo method for solving problems with a large number of dimensions.

- In this case typically the Monte Carlo method converge to the final value faster than the traditional numerical methods.

Pricing a Call Option

- It's worth to recast the pricing problem into a simple integral formulation in order to gain some insight into the general problem;
- So let's consider again the payoff of a simple plain vanilla option

$$
e^{-rT} \mathbb{E}^{\mathbb{Q}}[h(S_T)] = e^{-rT} \mathbb{E}^{\mathbb{Q}} \left[h \left(S_0 e^{\log(S_T/S_0)} \right) \right]
$$

- By a simple application of Ito's lemma is easy to demonstrate that the variable $X = \log(S_T/S_0)$ has a normal distribution with mean $m = (r - \frac{1}{2}\sigma^2)T$ and variance $s = \sigma^2 T$.
- So we can write

$$
C(S, t) = e^{-rT} \int_{-\infty}^{+\infty} \max[S_0 e^X - K, 0] e^{-\frac{(X-m)^2}{2s^2}} dX
$$

Pricing a Call Option

- It is possible to generate a normally distributed random variable $X = \Phi^{-1}(U; (r - \frac{1}{2}\sigma^2)T; \sigma^2 T)$ using the inverse transform method, where $\Phi^{-1}(\dots)$ is the inverse of the normal cumulative distribution function evaluated at U;

- U is a uniform $[0, 1]$ random variable.

$$U = \Phi[X; m, u], \quad u \to 1 \, when \, X \to +\infty, \quad u \to 0 \, when \, X \to -\infty$$

- From the previous relation we find (within a normalization factor)

$$du = \frac{d\Phi[X; m, u]}{dX} dX \Rightarrow dX = \frac{1}{e^{-\frac{(X-m)^2}{2s^2}}} du$$

- and ...

Pricing a Call Option

- ... finally we can write our integral in the form

$$C(S, t) = \int_0^1 f(u)\, du$$

where $f(u) = e^{-rT} \max[S_0 \exp(\Phi^{-1}(u; m, s)) - K, 0]$

Pricing a Call Option - The Python Code

```python
def f(u, S0, K, r, sigma, T):
    m      = (r - .5*sigma*sigma)*T
    s      = sigma*sqrt(T)
    f_u    = exp(-r*T) *
             np.maximum(S0*exp(scnorm.ppf(u, m, s))-K,0)
    return f_u


u      = rand(1000000)
f_u    = f(u,S0,K,r,sigma,T)


print mean(f_u)
```

Pricing a Call Option - The Integrand Function

Notebook

- **GitHub :**   polyhedron-gdl;
- **Notebooks :**   mcs_3;

Feynman–Kac formula

- The **Feynman–Kac formula** named after Richard Feynman and Mark Kac, establishes a link between parabolic partial differential equations (PDEs) and stochastic processes.

- It offers a method of solving certain PDEs by simulating random paths of a stochastic process. Conversely, an important class of expectations of random processes can be computed by deterministic methods.

- Consider the PDE

$$\frac{\partial u}{\partial t}(x, t) + \mu(x, t)\frac{\partial u}{\partial x}(x, t) + \tfrac{1}{2}\sigma^2(x, t)\frac{\partial^2 u}{\partial x^2}(x, t) - V(x, t)u(x, t) + f(x, t) = 0$$

subject to the terminal condition

$$u(x, T) = \psi(x)$$

Feynman–Kac formula

- Then the Feynman–Kac formula tells us that the solution can be written as a conditional expectation

$$u(x,t) = E^Q \left[\int_t^T e^{-\int_t^r V(X_\tau, \tau)\, d\tau} f(X_r, r) dr + e^{-\int_t^T V(X_\tau, \tau)\, d\tau} \psi(X_T) \,\middle|\, X_t = x \right]$$

under the probability measure $\mathbb{Q}$ such that X' is an Ito process driven by the equation

$$dX = \mu(X,t)\, dt + \sigma(X,t)\, dW^{\mathbb{Q}}$$

Valuing a derivative contract

- A derivative can be perfectly replicated by means of a self-financing dynamic portfolio whose value exactly matches all of the derivative flows in every state of the world. This approach shows that the values of the derivative (and of the portofolio) solves the following PDE

$$\frac{\partial f}{\partial t} + \frac{\partial f}{\partial S} rS + \frac{1}{2} \frac{\partial f}{\partial S} \sigma^2 S^2 = fr \tag{1}$$

  with the terminal condition at T that is the derivative's payoff

$$f(T, S(T)) = payoff$$

Valuing a derivative contract

- According to the Feynmann-Kac formula, if $f(t_0, S(t_0))$ solves the B-S PDE, then it is also solution of

$$f(t_0, S(t_0)) = \mathbb{E}\left[e^{-r(T-t_0)}f(T, S(T)|\mathcal{F}_{t_0}\right]$$

- i.e. it's the expected value of the discounted payoff in a probability measure where the evolution of the asset is

$$dS = rSdt + \sigma Sdw$$

- This probability measure is the **Risk Neutral Measure**

Valuing a derivative contract

- Since there exist such an equivalence, we can compute option prices by means of two numerical methods
- PDE: finite difference (explicit, implicit, crank-nicholson)suitable for optimal exercise derivatives;
- Integration
  - Quadrature Methods;
  - Monte Carlo Methods

Outline

Scenario Contruction

- There are several ways to construct scenario for pricing
- Constructing a path of the solution to the SDE at times t_i by exact advancement of the solution;
  - This method is only possible if we have an analytical expression for the solution of the stochastic differential equation
- Approximate numerical solution of the stochastic differential equation;
  - This is the method of choice if we cannot use the previous one; Just as in the case of ODE there are numerical techniques for discretizing and solving SDE.

Exact Solution Advancement

- Example: Log-normal process with constant drift and volatility

$$\text{SDE} \Rightarrow \frac{dS}{S} = \mu dt + \sigma dw$$

$$\text{SOLUTION} \Rightarrow S(T) = S(t_0) e^{(\mu - \frac{1}{2}\sigma^2)(T - t_0) + \sigma[w(T) - w(t_0)]}$$

- How to obtain a sequence of Wiener process?

$$w(t_i) = w(t_{i-1}) + \sqrt{t_i - t_{i-1}} Z \quad Z \sim N(0, 1)$$

Exact Solution Advancement

- Defining the outcomes of successive drawings of the random variable Z corresponding to the $j - th$ trajectory by Z_i^j, we get the following recursive expression for the $j - th$ trajectory of $S(t)$:

$$S^j(t_i) = S^j(t_{i-1}) exp \left[\left(\mu - \frac{1}{2}\sigma^2 \right) (t_i - t_{i-1}) + \sigma \sqrt{t_i - t_{i-1}} Z_i^j \right]$$

Numerical Integration of SDE

- The numerical integration of the SDE by finite difference is another way of generating scenarios for pricing;

- In the case of the numerical integration of ordinary differential equations by finite differences the numerical scheme introduces a discretization error that translates into the numerical solution differing from the exact solution by an amount proportional to a power of the time step.

- This amount is the truncation error of the numerical scheme.

Numerical Integration of SDE

- In the case of the numerical integration of SDE by finite differences the interpretation of the numerical error introduced by the discretization scheme is more complicated;

- Unlike the case of ODE where the only thing we are interested in computing is the solution itself, when dealing with SDE there are two aspects that interest us:
  - One aspect is the accuracy with which we compute the trajectories or paths of a realization of the solution
  - The other aspect is the accuracy with which we compute functions of the process such as expectations and moments.

Numerical Integration of SDE

- The order of accuracy with which a given scheme can approximate trajectories of the solution is not the same as the accuracy with which the same scheme can approximate expectations and moments of functions of the trajectories;

- The convergence of the numerically computed trajectories to the exact trajectories is called strong convergence and the order of the corresponding numerical scheme is called order of strong convergence;

- The convergence of numerically computed functions of the stochastic process to the exact values is called weak convergence and the related order is called order of weak convergence.

Numerical Integration of SDE

- We assume that the stock price S_t is driven by the stochastic differential equation (SDE)

$$dS(t) = \mu(S, t)dt + \sigma(S, t)dW_t \qquad (2)$$

  where W_t is, as usual, Brownian motion.

- We simulate S_t over the time interval $[0; T]$, which we assume to be is discretized as $0 = t_1 < t_2 < \cdots < t_m = T$, where the time increments are equally spaced with width dt.

- Equally-spaced time increments is primarily used for notational convenience, because it allows us to write $t_i - t_{i-1}$ as simply dt. All the results derived with equally-spaced increments are easily generalized to unequal spacing.

Numerical Integration of SDE

- **Euler Scheme**

- The simplest way to discretize the process in Equation (**??**) is to use Euler discretization

$$\text{EULER} \Rightarrow \hat{S}(t_{i+1}) = \hat{S}(t_i) + \mu[\hat{S}(t_i), t_i]\Delta t + \sigma[\hat{S}(t_i), t_i](w(t_{i+1}) - w(t_i))$$

- **Milshstein Scheme**

$$\text{MILSHSTEIN} \Rightarrow \text{EULER} + \frac{1}{2}\sigma[\hat{S}(t_i)]\frac{\partial\sigma[\hat{S}(t_i)]}{\partial S}\left[(w(t_{i+1}) - w(t_i))^2 - \Delta t\right]$$

This scheme is described in Glasserman and in Kloeden and Platen for general processes, and in Kahl and Jackel for stochastic volatility models. The scheme works for SDEs for which the coefficients $\mu(S_t)$ and $\sigma(S_t)$ depend only on S, and do not depend on t directly

Numerical Integration of SDE

- For an intuitive derivation, we will only look at GBM:

$$d \ln x_t = \left(\mu - \frac{1}{2}\sigma^2 \right) dt + \sigma dW_t$$

- The solution to the GBM SDE is

$$x_{t+\Delta t} = x_t \exp \left\{ \int_t^{t+\Delta t} \left(\mu - \frac{1}{2}\sigma^2 \right) dt + \int_t^{t+\Delta t} \sigma dW_u \right\}$$

$$\sim x_t \left(1 + \mu\Delta t - \frac{1}{2}\sigma^2\Delta t + \sigma\Delta W_t + \frac{1}{2}\sigma^2(\Delta W_t)^2 \right)$$

$$\Rightarrow x_{t+\Delta t} \sim x_t + a(x_t)\Delta t + b(x_t)\Delta w_t + \frac{1}{2}b(x_t)\frac{\partial b(x_t)}{\partial x_t}((\Delta W_t)^2 - \Delta t)$$

Notebook

- **GitHub :** polyhedron-gdl;
- **Notebooks :** mcs_sde_solution;
- **Code :** mcs_sde_solution.py;

The Brownian Bridge

- Assume you have a Wiener process defined by a set of time-indexed random variables $W(t_1), W(t_2), ..., W(t_n)$.

- How do you insert a random variable $W(t_k)$ where $t_i \leq t_k \leq t_{i+1}$ into the set in such a manner that the resulting set still constitutes a Wiener process?

- The answer is: with a Brownian Bridge!

- The Brownian Bridge is a sort of interpolation that allows you to introduce intermediate points in the trajectory of a Wiener process.

The Brownian Bridge

- Brownian Bridge Construction
- Given $W(t)$ and $W(t + \delta t_1 + \delta t_2)$ we want to find $W(t + \delta t_1)$;
- We assume that we can get the middle point by a weighted average of the two end points plus an independent normal random variable:

$$W(t + \delta t_1) = \alpha W(t) + \beta W(t + \delta t_1 + \delta t_2) + \gamma Z$$

where α, β and λ are constants to be determined and Z is a standard normal random variable.

The Brownian Bridge

- We have to satisfy the following conditions:

$$\begin{cases} cov[W(t+\Delta t_1), W(t)] = min(t+\Delta t_1, t) = t \\ cov[W(t+\Delta t_1), W(t+\Delta t_1+\Delta t_2)] = t+\Delta t_1 \\ var[W(t+\Delta t_1)] = t+\Delta t_1 \end{cases}$$

$$\begin{cases} \alpha + \beta = 1 \\ \alpha t + \beta(t+\Delta t_1+\Delta t_2) = t+\Delta t_1 \\ \alpha^2 t + 2\alpha\beta t + \beta^2(t+\Delta t_1+\Delta t_2) + \lambda^2 = t+\Delta t_1 \end{cases}$$

- which are equivalent to:

$$\begin{cases} \alpha = \frac{\Delta t_2}{\Delta t_1+\Delta t_2} \\ \beta = 1 - \alpha = \frac{\Delta t_1}{\Delta t_1+\Delta t_2} \\ \gamma = \sqrt{\Delta t_1 \alpha} \end{cases}$$

The Brownian Bridge

- We can use the brownian bridge to generate a Wiener path and then use the Wiener path to produce a trajectory of the process we are interested in;

- The simplest strategy for generating a Wiener path using the brownian bridge is to divide the time span of the trajectory into two equal parts and apply the brownian bridge construction to the middle point. We then repeat the procedure for the left and right sides of the time interval.

The Brownian Bridge

- Notice that as we fill in the Wiener path, the additional variance of the normal components we add has decreasing value;

- Of course the total variance of all the Wiener increments does not depend on how we construct the path, however the fact that in the brownian bridge approach we use random variables that are multiplied by a factor of decreasing magnitude means that the importance of those variables also decreases as we fill in the path;

- The dimension of the random variables with larger variance need to be sampled more efficiently than the dimension with smaller variance;

Notebook

- **GitHub :** polyhedron-gdl;
- **Code :** mcs_brownian_bridge.py;

Outline

Variance Reduction Methods

- In this section we briefly discuss techniques for improving on the speed and efficiency of a simulazion, usually called *variance reduction techniques*;

- If we do nothing about efficiency, the number of MC replications we need to achieve acceptable pricing acccuracy may be surprisingly large;

- As a result in many cases variance reduction techiques are a practical requirement;

- From a general point of view these methods are based on two principal strategies for reducing variance:
  - Taking advantage of tractable features of a model to adjust or correct simulation output
  - Reducing the variability in simulation input

Variance Reduction Methods

- From the first section we remember that the variance of the estimator is

$$var\left(\tilde{I}_n\right) = \frac{var(f(U_i))}{n}$$

- So, the standard error of the sample mean is the standard deviation or

$$SE\left(\tilde{I}_n\right) = \frac{\sigma_f}{\sqrt{n}}$$

where $\sigma_f^2 = var(f(U_i))$

Variance Reduction Methods

- The most commonly used strategies for variance reduction are the following:
  - **Antithetic variates**
  - **Moment Matching**
  - Control variates
  - **Stratified Sampling**
  - Importance Sampling
  - Low-discrepancy sequences

Variance Reduction Methods - Antithetic Variates

- In this case we construc the estimator by using two brownian trajectories that are mirror images of each other;
- This causes cancellation of dispersion;
- This method tends to reduce the variance modestly but it is extremely easy to implement and as a result very commonly used;
- For the antithetic method to work we need V^+ and V^- to be negatively correlated;
- this will happen if the payoff function is a monotonic function of Z;

Variance Reduction Methods - Antithetic Variates

- To apply the antithetic variate technique, we generate standard normal random numbers Z and define two set of samples of the undelying price

$$S_T^+ = S_0 e^{(r-\sigma^2/2)T + \sigma\sqrt{T}Z} \quad S_T^- = S_0 e^{(r-\sigma^2/2)T + \sigma\sqrt{T}(-Z)}$$

- Similarly we define two sets of discounted payoff samples ...

$$V_T^+ = \max[S^+(T) - K, 0] \quad V_T^- = \max[S^-(T) - K, 0]$$

- ... and at last we construct our mean estimator by averaging these samples

$$\bar{V}_0 = \frac{1}{n} \sum_{j=1}^{n} \frac{1}{2} \left(V_j^+ + V_j^- \right)$$

Variance Reduction Methods - Moment Matching

- Let $z_i, i = 1, ..., n$, denote an independent standard normal random vector used to drive a simulation.

- The sample moments will not exactly match those of the standard normal. The idea of moment matching is to transform the z_i to match a finite number of the moments of the underlying population.

- For example, the first and second moment of the normal random number can be matched by defining

$$\tilde{z}_i = (z_i - \tilde{z})\frac{\sigma_z}{s_z} + \mu_z, i = 1, ...n \qquad (3)$$

where $\tilde{z}$ is the sample mean of the z_i and σ_z is the population standard deviation, s_z is the sample standard deviation of z_i, and μ_z s the population mean.

Notebook

- **GitHub :** polyhedron-gdl;
- **Notebook :** n03_mcs;

Outline

Choleski Decomposition

- The **Choleski Decomposition** makes an appearance in Monte Carlo Methods where it is used to simulate systems with correlated variables.
- Cholesky decomposition is applied to the correlation matrix, providing a lower triangular matrix A, which when applied to a vector of uncorrelated samples, u, produces the covariance vector of the system. Thus it is highly relevant for quantitative trading.
- The standard procedure for generating a set of correlated normal random variables is through a linear combination of uncorrelated normal random variables;
- Assume we have a set of n independent standard normal random variables Z and we want to build a set of n correlated standard normals Z' with correlation matrix Σ

$$Z' = AZ, \qquad AA^t = \Sigma$$

Choleski Decomposition

- We can find a solution for A in the form of a triangular matrix

$$\begin{pmatrix} A_{11} & 0 & \ldots & 0 \\ A_{21} & A_{22} & \ldots & 0 \\ \vdots & \vdots & \ddots & \ldots \\ A_{n1} & A_{n2} & \ldots & A_{nn} \end{pmatrix}$$

- **diagonal elements**

$$a_{ii} = \sqrt{\Sigma_{ii} - \sum_{k=1}^{i-1} a_{ik}^2}$$

- **off-diagonal elements**

$$a_{ij} = \frac{1}{a_{ii}} \left(\Sigma_{ij} - \sum_{k=1}^{i-1} a_{ik} a_{jk} \right)$$

Choleski Decomposition

- For example, for a two-dimension random vector we have simply

$$A = \begin{pmatrix} \sigma_1 & 0 \\ \sigma_2 \rho & \sigma_2 \sqrt{1 - \rho^2} \end{pmatrix}$$

- say one needs to generate two correlated normal variables x_1 and x_2
- All one needs to do is to generate two uncorrelated Gaussian random variables z_1 and z_2 and set

$$x_1 = z_1$$

$$x_2 = \rho z_1 + \sqrt{1 - \rho^2} z_2$$

Notebook

- **GitHub :** polyhedron-gdl;
- **Notebook :** n05_mcs_multi_asset_path;

Outline

CIR Model

- In this section, we consider the stochastic short rate model MCIR85 of Cox- Ingersoll-Ross which is given by the SDE:

$$dr_t = \kappa_r(\theta_r - r_t)dt + \sigma_r\sqrt{r_t}dZ_t \qquad (4)$$

- To simulate the short rate model, it has to be discretized. To this end, we divide the given time interval $[0, T]$ in equidistant sub-intervals of length t such that now $t \in \{0, \Delta t, 2\Delta t, \dots, T\}$, i.e. there are $M + 1$ points in time with $M = T/t$.

- The exact transition law of the square-root diffusion is known. Consider the general square- root diffusion process

$$dx_t = \kappa(\theta - x_t)dt + \sigma\sqrt{x_t}dZ_t \qquad (5)$$

CIR Model

- It can be show that x_t, given x_s with $s = t - \Delta t$, is distributed according to

$$x_t = \frac{\sigma^2(1 - e^{-\kappa\Delta t})}{4\kappa} \chi_d'^2 \left(\frac{4^{-\kappa\Delta t}}{\sigma^2(1 - e^{-\kappa\Delta t})} x_s \right)$$

where $\chi_d'^2$ denotes a non-central chi-squared random variable with

$$d = \frac{4\theta\kappa}{\sigma^2}$$

degrees of freedom and non-centrality parameter

$$l = \frac{4^{-\kappa\Delta t}}{\sigma^2(1 - e^{-\kappa\Delta t})} x_s$$

CIR Model

- For implementation purposes, it may be convenient to sample a chi-squared random variable χ_d^2 instead of a non-central chi-squared one, $\chi_d'^2$.

- If $d > 1$, the following relationship holds true

$$\chi_d'^2(l) = (z + \sqrt{l})^2 + \chi_{d-1}^2$$

where z is an independent standard normally distributed random variable.

- Similarly, if $d \leq 1$, one has

$$\chi_d'^2(l) = \chi_{d+2N}^2$$

where N is now a Poisson-distributed random variable with intensity $l/2$. For an algorithmic representation of this simulation scheme refer to Glasserman, p. 124.

CIR Model: Pricing ZCB

- A MC estimator for the value of the ZCB at t is derived as follows.
- Consider a certain path i of the I simulated paths for the short rate process with time grid $t \in \{0, \Delta t, 2\Delta t, \ldots, T\}$.
- We discount the terminal value of the ZCB, i.e. 1, step-by-step backward. For $t < T$ and $s = t - \Delta t$ we have

$$B_{s,i} = B_{t,i} e^{-\frac{r_t + r_s}{2} \Delta t}$$

- The MC estimator of the ZCB value at t is

$$B_t^{MC} = \frac{1}{I} \sum_{i=1}^{I} B_{t,i}$$

CIR Model: Pricing ZCB

- The present value of the ZCB in the CIR model takes the form:

$$B_0(T) = b_1(T)e^{-b_2(T)r_0}$$

where

$$b_1(T) = \left[\frac{2\gamma \, exp((\kappa_r + \gamma)T/2)}{2\gamma + (\kappa_r + \gamma)(e^{\gamma T} - 1)} \right]^{\frac{2\kappa_r\theta_r}{\sigma_r^2}}$$

$$b_2(T) = \frac{2(e^{\gamma T} - 1)}{2\gamma + (\kappa_r + \gamma)(e^{\gamma T} - 1)}$$

$$\gamma = \sqrt{\kappa_r^2 + 2\sigma_r^2}$$

Notebook

- **GitHub :**   polyhedron-gdl;
- **Notebook :**   n06_mcs_cir;

The Heston Model

- Stochastic volatility models are those in which the variance of a stochastic process is itself randomly distributed.

- The models assumes that the underlying security's volatility is a random process, governed by state variables such as the price level of the underlying security, the tendency of volatility to revert to some long-run mean value, and the variance of the volatility process itself, among others.

- Stochastic volatility models are one approach to resolve a shortcoming of the Black–Scholes model.

- In particular this model cannot explain long-observed features of the implied volatility surface such as volatility smile and skew, which indicate that implied volatility does tend to vary with respect to strike price and expiry.

The Heston Model

By assuming that the volatility of the underlying price is a stochastic process rather than a constant, it becomes possible to model derivatives more accurately.

- Heston model
- CEV model
- SABR volatility model
- GARCH model

The Heston Model

- In this section we are going to consider the stochastic volatility model MH93 with constant short rate.

- This section values European call and put options in this model by MCS.

- As for the ZCB values, we also have available a semi-analytical pricing formula which generates natural benchmark values against which to compare the MCS estimates.

The Heston Model

- The basic Heston model assumes that S_t, the price of the asset, is determined by a stochastic process:

$$dS_t = \mu S_t\, dt + \sqrt{\nu_t} S_t\, dW_t^S$$

where ν_t, the instantaneous variance, is a CIR process:

$$d\nu_t = \kappa(\theta - \nu_t)\, dt + \xi\sqrt{\nu_t}\, dW_t^\nu$$

and dW_t^S, dW_t^ν are Wiener process with correlation ρ, or equivalently, with covariance ρdt.

The Heston Model

The parameters in the above equations represent the following:

- μ is the rate of return of the asset.

- θ is the *long variance*, or long run average price variance; as t tends to infinity, the expected value of ν_t tends to θ.

- κ is the rate at which ν_t reverts to θ.

- ξ is the volatility of the volatility, or *vol of vol*, and determines the variance of ν_t.

If the parameters obey the following condition (known as the Feller condition) then the process ν_t is strictly positive

$$2\kappa\theta > \xi^2$$

The Heston Model

- The correlation introduces a new problem dimension into the discretization for simulation purposes.

- To avoid problems arising from correlating normally distributed increments (of S) with chi-squared distributed increments (of v), we will in the following only consider Euler schemes for both the S and v process.

- This has the advantage that the increments of v become normally distributed as well and can therefore be easily correlated with the increments of S.

The Heston Model

- we consider two discretization schemes for S and seven discretization schemes for v.

- For S we have the simple Euler discretization scheme (with $s = t - \Delta t$)

$$S_t = S_s \left(e^{r\Delta t} + \sqrt{v_t}\sqrt{\Delta t}z_t^1 \right)$$

As an alternative we consider the exact log Euler scheme

$$S_t = S_s e^{(r - v_t/2)\Delta t + \sqrt{v_t}\sqrt{\Delta t}z_t^1}$$

This one is obtained by considering the dynamics of $\log S_t$ and applying Ito's lemma to it.

The Heston Model

- These schemes can be combined with any of the following Euler schemes for the square-root diffusion ($x^+ = \max[0, x]$):

- **Full Truncation**

$$\tilde{x}_t = \tilde{x}_s + \kappa(\theta - \tilde{x}_s^+)\Delta t + \sigma\sqrt{\tilde{x}_s^+}\sqrt{\Delta t}z_t, \quad x_t = \tilde{x}_t^+$$

- **Partial Truncation**

$$\tilde{x}_t = \tilde{x}_s + \kappa(\theta - \tilde{x}_s)\Delta t + \sigma\sqrt{\tilde{x}_s^+}\sqrt{\Delta t}z_t, \quad x_t = \tilde{x}_t^+$$

- **Truncation**

$$x_t = \max\left[0, \tilde{x}_s + \kappa(\theta - \tilde{x}_s)\Delta t + \sigma\sqrt{\tilde{x}_s}\sqrt{\Delta t}z_t\right]$$

- **Reflection**

$$\tilde{x}_t = |\tilde{x}_s| + \kappa(\theta - |\tilde{x}_s|)\Delta t + \sigma\sqrt{|\tilde{x}_s|}\sqrt{\Delta t}z_t, \quad x_t = |\tilde{x}_t|$$

The Heston Model

- **Hingham-Mao**

$$\tilde{x}_t = \tilde{x}_s + \kappa(\theta - \tilde{x}_s)\Delta t + \sigma\sqrt{|\tilde{x}_s|}\sqrt{\Delta t}z_t, \quad x_t = |\tilde{x}_t|$$

- **Simple Reflection**

$$\tilde{x}_t = \left| \tilde{x}_s + \kappa(\theta - \tilde{x}_s)\Delta t + \sigma\sqrt{\tilde{x}_s}\sqrt{\Delta t}z_t \right|$$

- **Absorption**

$$\tilde{x}_t = \tilde{x}_s^+ + \kappa(\theta - \tilde{x}_s^+)\Delta t + \sigma\sqrt{\tilde{x}_s^+}\sqrt{\Delta t}z_t, \quad x_t = \tilde{x}_t^+$$

- In the literature there are a lot of tests and numerical studies available that compare efficiency and precision of different discretization schemes.

Notebook

- **GitHub :**   polyhedron-gdl;
- **Notebook :**   n07_mcs_heston;